



```
In [46]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [47]: df = pd.read_csv('uber.csv')
```

```
In [48]: df.head()
```

```
Out[48]:
```

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	24238194	2015-05-07 19:52:06	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.712874	-73.999817	40.712874	1
1	27835199	2009-07-17 20:04:56	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.712874	-73.994355	40.712874	1
2	44984355	2009-08-24 21:45:00	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.712874	-74.005043	40.712874	1
3	25894730	2009-06-26 08:22:21	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.712874	-73.976124	40.712874	1
4	17610152	2014-08-28 17:47:00	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.712874	-73.925023	40.712874	1

```
In [49]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Unnamed: 0        200000 non-null  int64  
 1   key              200000 non-null  object 
 2   fare_amount      200000 non-null  float64 
 3   pickup_datetime  200000 non-null  object 
 4   pickup_longitude 200000 non-null  float64 
 5   pickup_latitude  200000 non-null  float64 
 6   dropoff_longitude 199999 non-null  float64 
 7   dropoff_latitude 199999 non-null  float64 
 8   passenger_count  200000 non-null  int64  
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

```
In [50]: df.columns
```

```
Out[50]: Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
               'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
               'dropoff_latitude', 'passenger_count'],
               dtype='object')
```

```
In [62]: df = df.drop(['Unnamed: 0', 'key'], axis=1)
```

```
In [63]: df.shape
```

```
Out[63]: (200000, 7)
```

```
In [64]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 7 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   fare_amount      200000 non-null   float64
 1   pickup_datetime  200000 non-null   object  
 2   pickup_longitude 200000 non-null   float64
 3   pickup_latitude   200000 non-null   float64
 4   dropoff_longitude 199999 non-null   float64
 5   dropoff_latitude  199999 non-null   float64
 6   passenger_count   200000 non-null   int64  
dtypes: float64(5), int64(1), object(1)
memory usage: 10.7+ MB
```

```
In [65]: df.describe()
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	11.359955	-72.527638	39.935885	-72.525292	40.734796
std	9.901776	11.437787	7.720539	13.117408	16.44421482
min	-52.000000	-1340.648410	-74.015515	-3356.666300	1644.421482
25%	6.000000	-73.992065	40.734796	-73.991407	1153.572603
50%	8.500000	-73.981823	40.752592	-73.980093	57.418457
75%	12.500000	-73.967153	40.767158	-73.963659	1644.421482
max	499.000000	57.418457	1644.421482	1153.572603	1644.421482

```
In [66]: df.dtypes
```

```
fare_amount          float64
pickup_datetime     object  
pickup_longitude    float64
pickup_latitude     float64
dropoff_longitude   float64
dropoff_latitude    float64
passenger_count     int64  
dtype: object
```

```
In [67]: df.isnull().sum()
```

```
Out[67]: fare_amount      0
         pickup_datetime    0
         pickup_longitude    0
         pickup_latitude     0
         dropoff_longitude   1
         dropoff_latitude    1
         passenger_count     0
dtype: int64
```

```
In [68]: df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(), inplace=True)
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].mean(), inplace=True)
```

C:\Users\Rohan\AppData\Local\Temp\ipykernel_1608\3614752822.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['dropoff_latitude'].fillna(value=df['dropoff_latitude'].mean(), inplace=True)
```

C:\Users\Rohan\AppData\Local\Temp\ipykernel_1608\3614752822.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df['dropoff_longitude'].fillna(value=df['dropoff_longitude'].mean(), inplace=True)
```

```
In [69]: df.isnull().sum()
```

```
Out[69]: fare_amount      0
         pickup_datetime    0
         pickup_longitude    0
         pickup_latitude     0
         dropoff_longitude   0
         dropoff_latitude    0
         passenger_count     0
dtype: int64
```

```
In [70]: df.pickup_datetime = pd.to_datetime(df.pickup_datetime, errors='coerce')
```

```
In [72]: df = df.assign(hour = df.pickup_datetime.dt.hour,
                      day = df.pickup_datetime.dt.day,
                      month = df.pickup_datetime.dt.month,
                      year = df.pickup_datetime.dt.year,
                      dayofweek = df.pickup_datetime.dt.dayofweek)
```

```
In [73]: df.head()
```

```
Out[73]:   fare_amount  pickup_datetime  pickup_longitude  pickup_latitude  dropoff_longitude
0          7.5  2015-05-07
              19:52:06+00:00           -73.999817        40.738354         -73.0
1          7.7  2009-07-17
              20:04:56+00:00           -73.994355        40.728225         -73.0
2         12.9  2009-08-24
              21:45:00+00:00           -74.005043        40.740770         -73.0
3          5.3  2009-06-26
              08:22:21+00:00           -73.976124        40.790844         -73.0
4         16.0  2014-08-28
              17:47:00+00:00           -73.925023        40.744085         -73.0
```


```
In [75]: df = df.drop('pickup_datetime', axis=1)
```

```
In [76]: df.head()
```

```
Out[76]:   fare_amount  pickup_longitude  pickup_latitude  dropoff_longitude  dropoff_latitude
0          7.5           -73.999817        40.738354         -73.999512        40.7
1          7.7           -73.994355        40.728225         -73.994710        40.7
2         12.9           -74.005043        40.740770         -73.962565        40.7
3          5.3           -73.976124        40.790844         -73.965316        40.8
4         16.0           -73.925023        40.744085         -73.973082        40.7
```

```
In [77]: df.plot(kind = 'box', subplots=True, layout=(7,2), figsize=(15,20))
```

```
Out[77]: fare_amount      Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude    Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude       Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude    Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude      Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count      Axes(0.547727,0.560732;0.352273x0.0939024)
hour                  Axes(0.125,0.448049;0.352273x0.0939024)
day                   Axes(0.547727,0.448049;0.352273x0.0939024)
month                 Axes(0.125,0.335366;0.352273x0.0939024)
year                  Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek               Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object
```




```
In [78]: #Using the interquartile range to fill the values
def remove_outlier(df1, col):
    Q1 = df1[col].quantile(0.25)
    Q3 = df1[col].quantile(0.75)
    IQR = Q3 - Q1
    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR
    df1[col] = np.clip(df1[col], lower, upper)
    return df1

def treat_outliers_all(df1, col_list):
    for c in col_list:
```

```
    df1 = remove_outlier(df, c)
return df1
```

```
In [79]: df.plot(kind='box', subplots=True, layout=(7,2), figsize=(15,20))
```

```
Out[79]: fare_amount           Axes(0.125,0.786098;0.352273x0.0939024)
pickup_longitude        Axes(0.547727,0.786098;0.352273x0.0939024)
pickup_latitude          Axes(0.125,0.673415;0.352273x0.0939024)
dropoff_longitude        Axes(0.547727,0.673415;0.352273x0.0939024)
dropoff_latitude          Axes(0.125,0.560732;0.352273x0.0939024)
passenger_count          Axes(0.547727,0.560732;0.352273x0.0939024)
hour                      Axes(0.125,0.448049;0.352273x0.0939024)
day                       Axes(0.547727,0.448049;0.352273x0.0939024)
month                     Axes(0.125,0.335366;0.352273x0.0939024)
year                      Axes(0.547727,0.335366;0.352273x0.0939024)
dayofweek                 Axes(0.125,0.222683;0.352273x0.0939024)
dtype: object
```



```
In [85]: import haversine as hs

travel_dist = []

# Ensure valid coordinates
df = df.loc[
    (df['pickup_latitude'].between(-90, 90)) &
    (df['dropoff_latitude'].between(-90, 90)) &
    (df['pickup_longitude'].between(-180, 180)) &
    (df['dropoff_longitude'].between(-180, 180))
].copy()
```

```

for pos in range(len(df)):
    lat1, lon1 = df['pickup_latitude'].iloc[pos], df['pickup_longitude'].iloc[pos]
    lat2, lon2 = df['dropoff_latitude'].iloc[pos], df['dropoff_longitude'].iloc[pos]
    loc1 = (lat1, lon1)
    loc2 = (lat2, lon2)
    dist = hs.haversine(loc1, loc2) # distance in kilometers
    travel_dist.append(dist)

df['dist_travel_km'] = travel_dist
df.head()

```

Out[85]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	-73.999817	40.738354	-73.999512	40.7
1	7.7	-73.994355	40.728225	-73.994710	40.7
2	12.9	-74.005043	40.740770	-73.962565	40.7
3	5.3	-73.976124	40.790844	-73.965316	40.8
4	16.0	-73.925023	40.744085	-73.973082	40.7

In [87]:

```
# Uber doesn't travel over 130kms so minimize distance
df = df.loc[(df.dist_travel_km >= 1) | (df.dist_travel_km <= 130)]
print("Remaining observations in the dataset:", df.shape)
```

Remaining observations in the dataset: (199988, 12)

In [88]:

```
#Finding incorrect latitude (Less than or greater than 90) and longitude (greater than 180)
incorrect_coordinates = df.loc[(df.pickup_latitude > 90) | (df.pickup_latitude < -90) |
                                 (df.dropoff_latitude > 90) | (df.dropoff_latitude < -90) |
                                 (df.pickup_longitude > 180) | (df.pickup_longitude < -180) |
                                 (df.dropoff_longitude > 90) | (df.dropoff_longitude < -90)]
```

In [90]:

```
df.drop(incorrect_coordinates, inplace=True, errors='ignore')
```

In [91]:

```
df.head()
```

Out[91]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
0	7.5	-73.999817	40.738354	-73.999512	40.7
1	7.7	-73.994355	40.728225	-73.994710	40.7
2	12.9	-74.005043	40.740770	-73.962565	40.7
3	5.3	-73.976124	40.790844	-73.965316	40.8
4	16.0	-73.925023	40.744085	-73.973082	40.7


In [92]:

```
df.isnull().sum()
```

```
Out[92]: fare_amount      0  
pickup_longitude      0  
pickup_latitude       0  
dropoff_longitude     0  
dropoff_latitude      0  
passenger_count       0  
hour                   0  
day                     0  
month                  0  
year                   0  
dayofweek              0  
dist_travel_km        0  
dtype: int64
```

```
In [93]: sns.heatmap(df.isnull())
```

```
Out[93]: <Axes: >
```



```
In [94]: corr = df.corr()
```


```
In [95]: corr
```

Out[95]:

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude
fare_amount	1.000000	0.011635	-0.011561	0.01
pickup_longitude	0.011635	1.000000	-0.979048	0.94
pickup_latitude	-0.011561	-0.979048	1.000000	-0.93
dropoff_longitude	0.011870	0.949698	-0.936690	1.00
dropoff_latitude	-0.012258	-0.936642	0.958143	-0.97
passenger_count	0.010159	-0.000756	-0.000523	-0.00
hour	-0.021495	0.001393	-0.001346	0.00
day	0.001395	0.005368	-0.008814	0.00
month	0.023795	-0.003574	0.004718	-0.00
year	0.118339	0.009736	-0.012036	0.01
dayofweek	0.007492	0.002312	-0.004148	0.00
dist_travel_km	0.024723	0.163547	-0.142200	0.14

In [96]: `fig, axis = plt.subplots(figsize=(10,6))
sns.heatmap(df.corr(), annot=True)`

Out[96]: <Axes: >



```
In [100... x = df[['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_l  
y = df['fare_amount']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.33)
```

```
In [ ]: # Linear Regression
from sklearn.linear_model import LinearRegression
regression = LinearRegression()
regression.fit(X_train, y_train)
```

```
Out[ ]: ▾ LinearRegression ⓘ ?
```

Parameters		
fit_intercept	True	
copy_X	True	
tol	1e-06	
n_jobs	None	
positive	False	

```
In [102... regression.intercept_
```

```
Out[102... np.float64(-1308.5763838650173)
```

```
In [103... regression.coef_
```

```
Out[103... array([ 1.82048902e-02,  6.29490061e-02, -2.44374099e-02, -8.05621034e-02,
                  6.22658606e-02, -3.02383898e-02,  4.64851748e-03,  1.05018214e-01,
                  6.55984557e-01,  2.34767432e-02,  5.40340472e-04])
```

```
In [105... prediction = regression.predict(X_test)
prediction
```

```
Out[105... array([12.42117782, 11.66326224, 11.46637051, ..., 10.86121541,
                  11.32915104, 12.33336062], shape=(65997,))
```

```
In [106... y_test
```

```
Out[106...]:
```

52153	19.5
47578	8.9
39397	8.0
35679	5.5
63534	8.5
	...
66104	6.1
41390	12.0
76456	8.5
147242	11.7
19361	18.0

```
Name: fare_amount, Length: 65997, dtype: float64
```

```
In [107...]: #R2, MSE, RMSE  
from sklearn.metrics import r2_score  
r2_score(y_test, prediction)
```

```
Out[107...]: 0.016433455965117583
```

```
In [108...]: from sklearn.metrics import mean_squared_error  
MSE = mean_squared_error(y_test, prediction)  
MSE
```

```
Out[108...]: 96.48994462606969
```

```
In [112...]: RMSE = np.sqrt(MSE)  
RMSE
```

```
Out[112...]: np.float64(9.822929533803533)
```